


Cross-attention paired-location endpoint replication is estimand-dependent



Each line is one independently trained checkpoint; V+M co-located totals are secondary, not source maps. Points average 128 held-out trials at frame 5; bars are trial-bootstrap 95% CIs, not training-population intervals.